

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

RocketMQ Operator

24 Nov, 2018

谐云科技 魏欢

阿里巴巴中间件

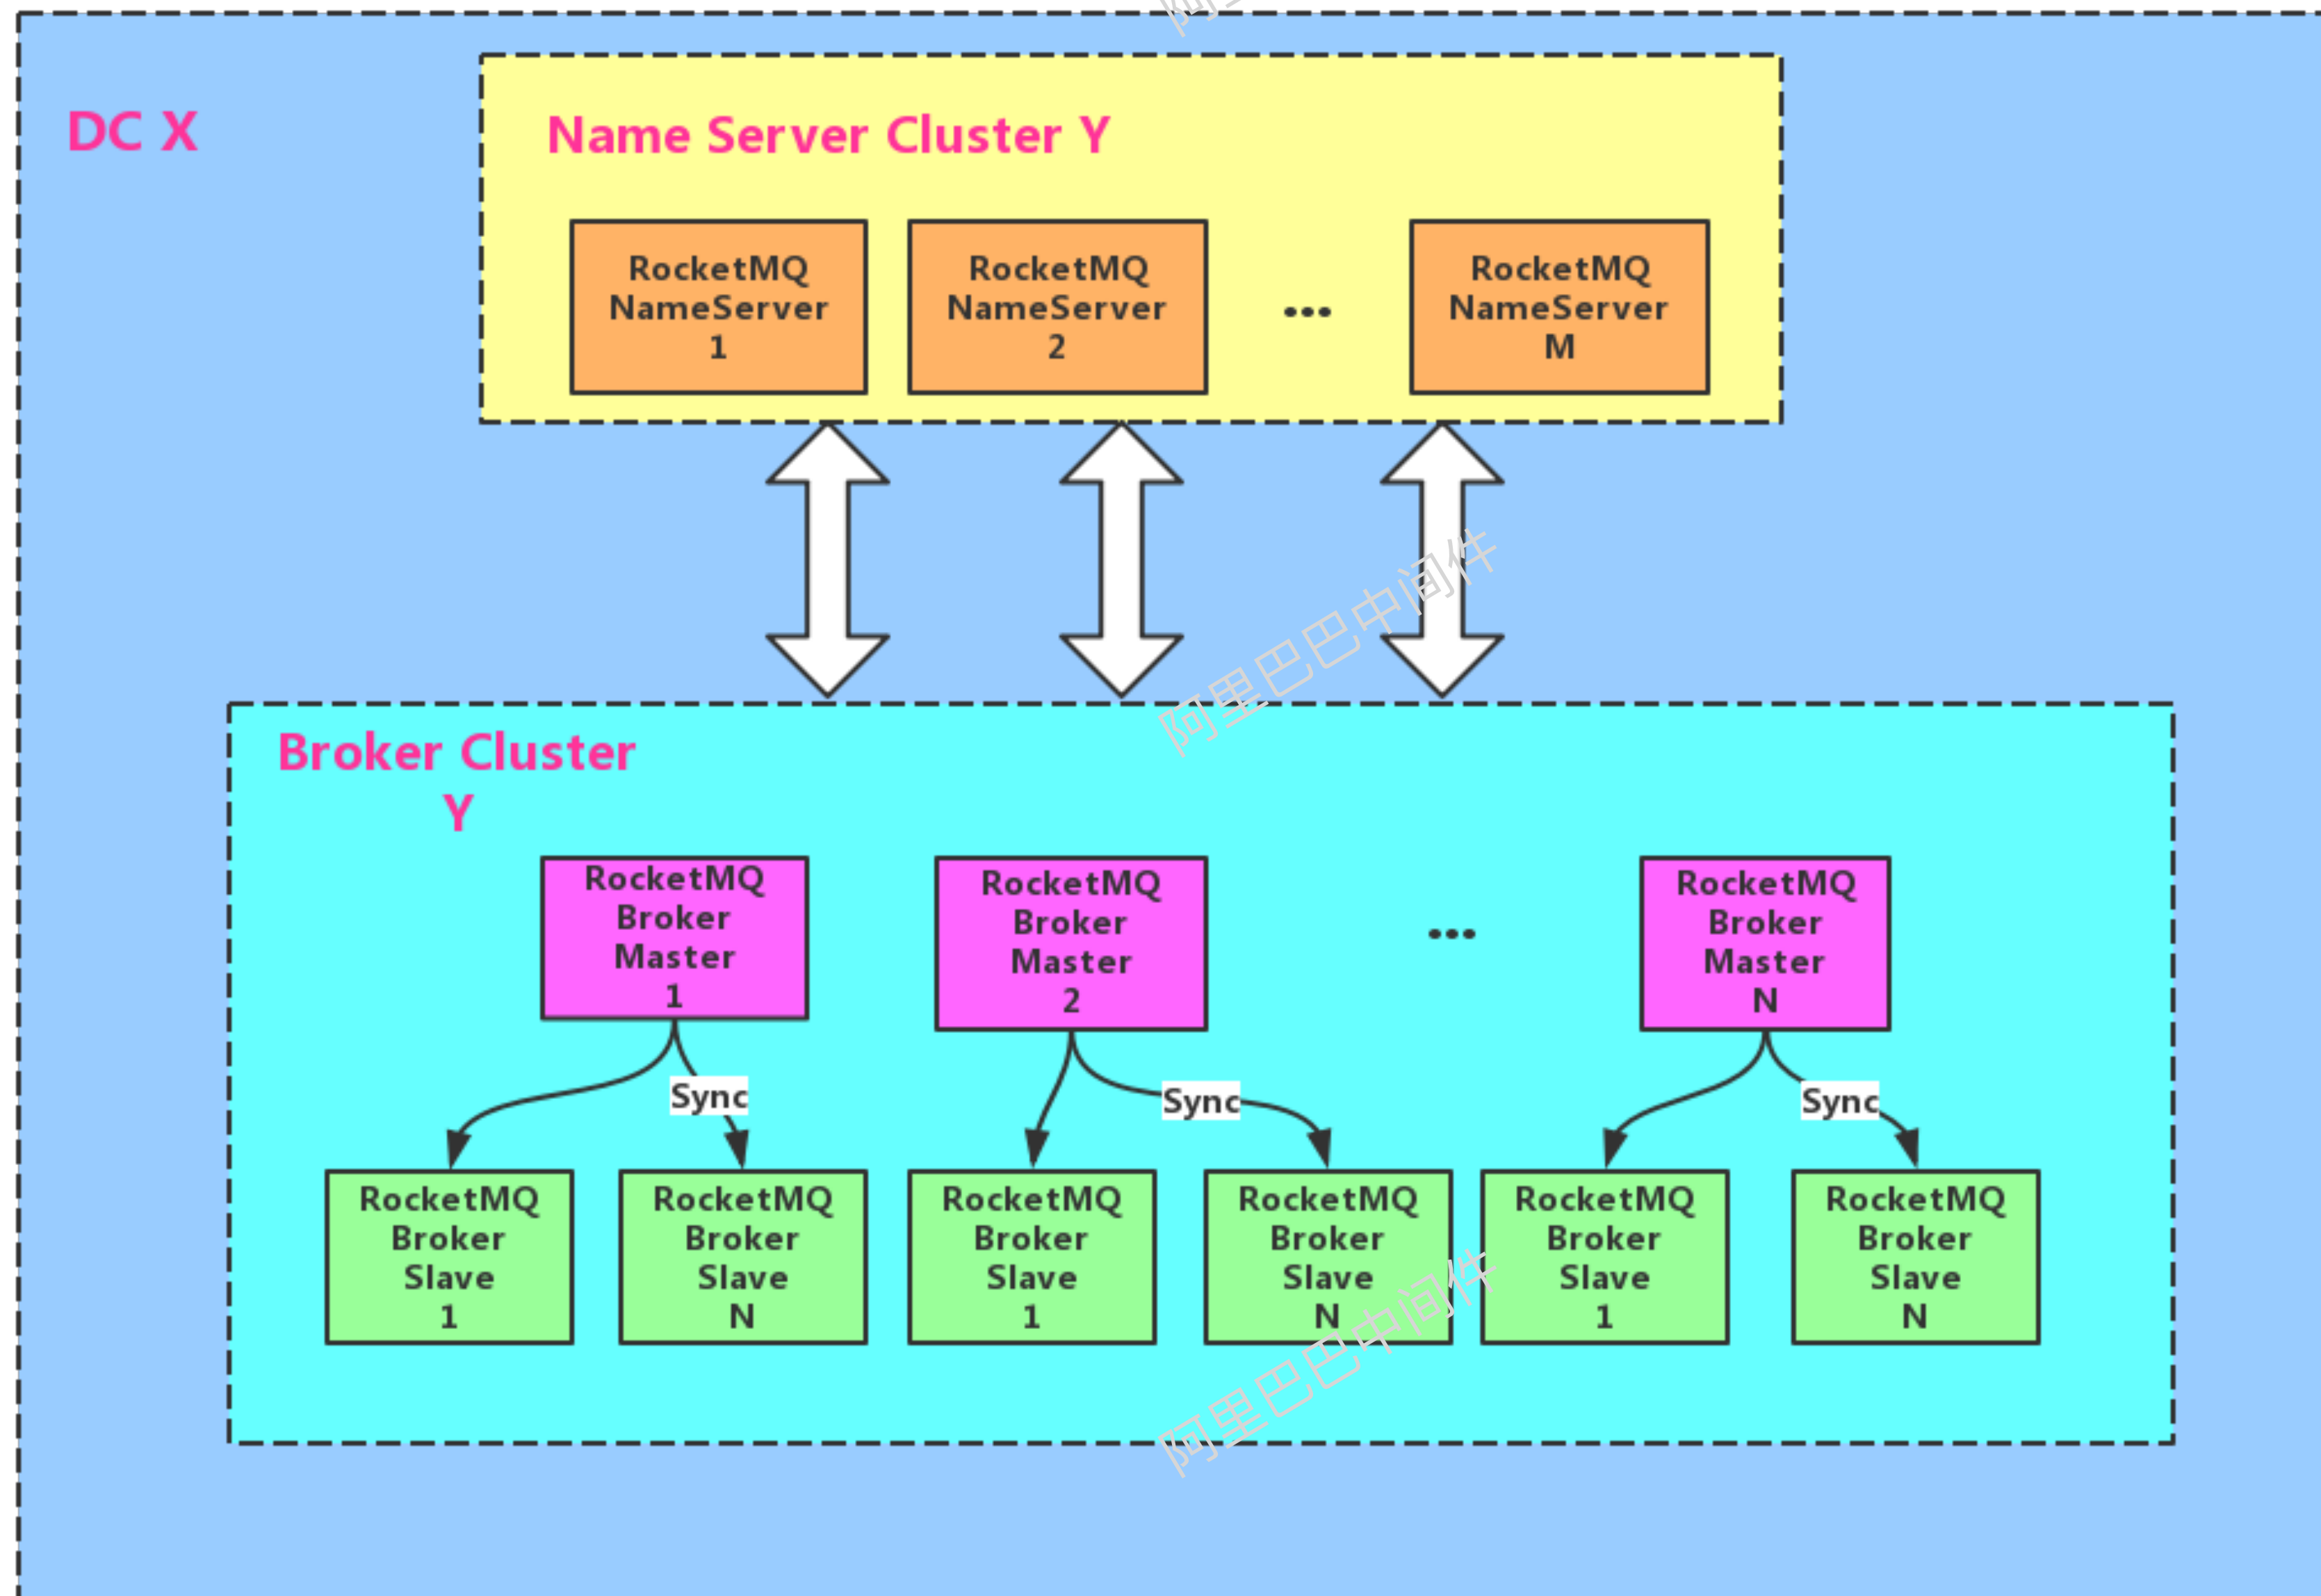
阿里巴巴中间件

阿里巴巴中间件

Agenda

- What is RocketMQ Operator
- How to design an Operator
- How to write RocketMQ Operator
- RocketMQ Operator Tutorial
- RocketMQ Operator Improvements Plan

Sample: RocketMQ Clusters In DC X



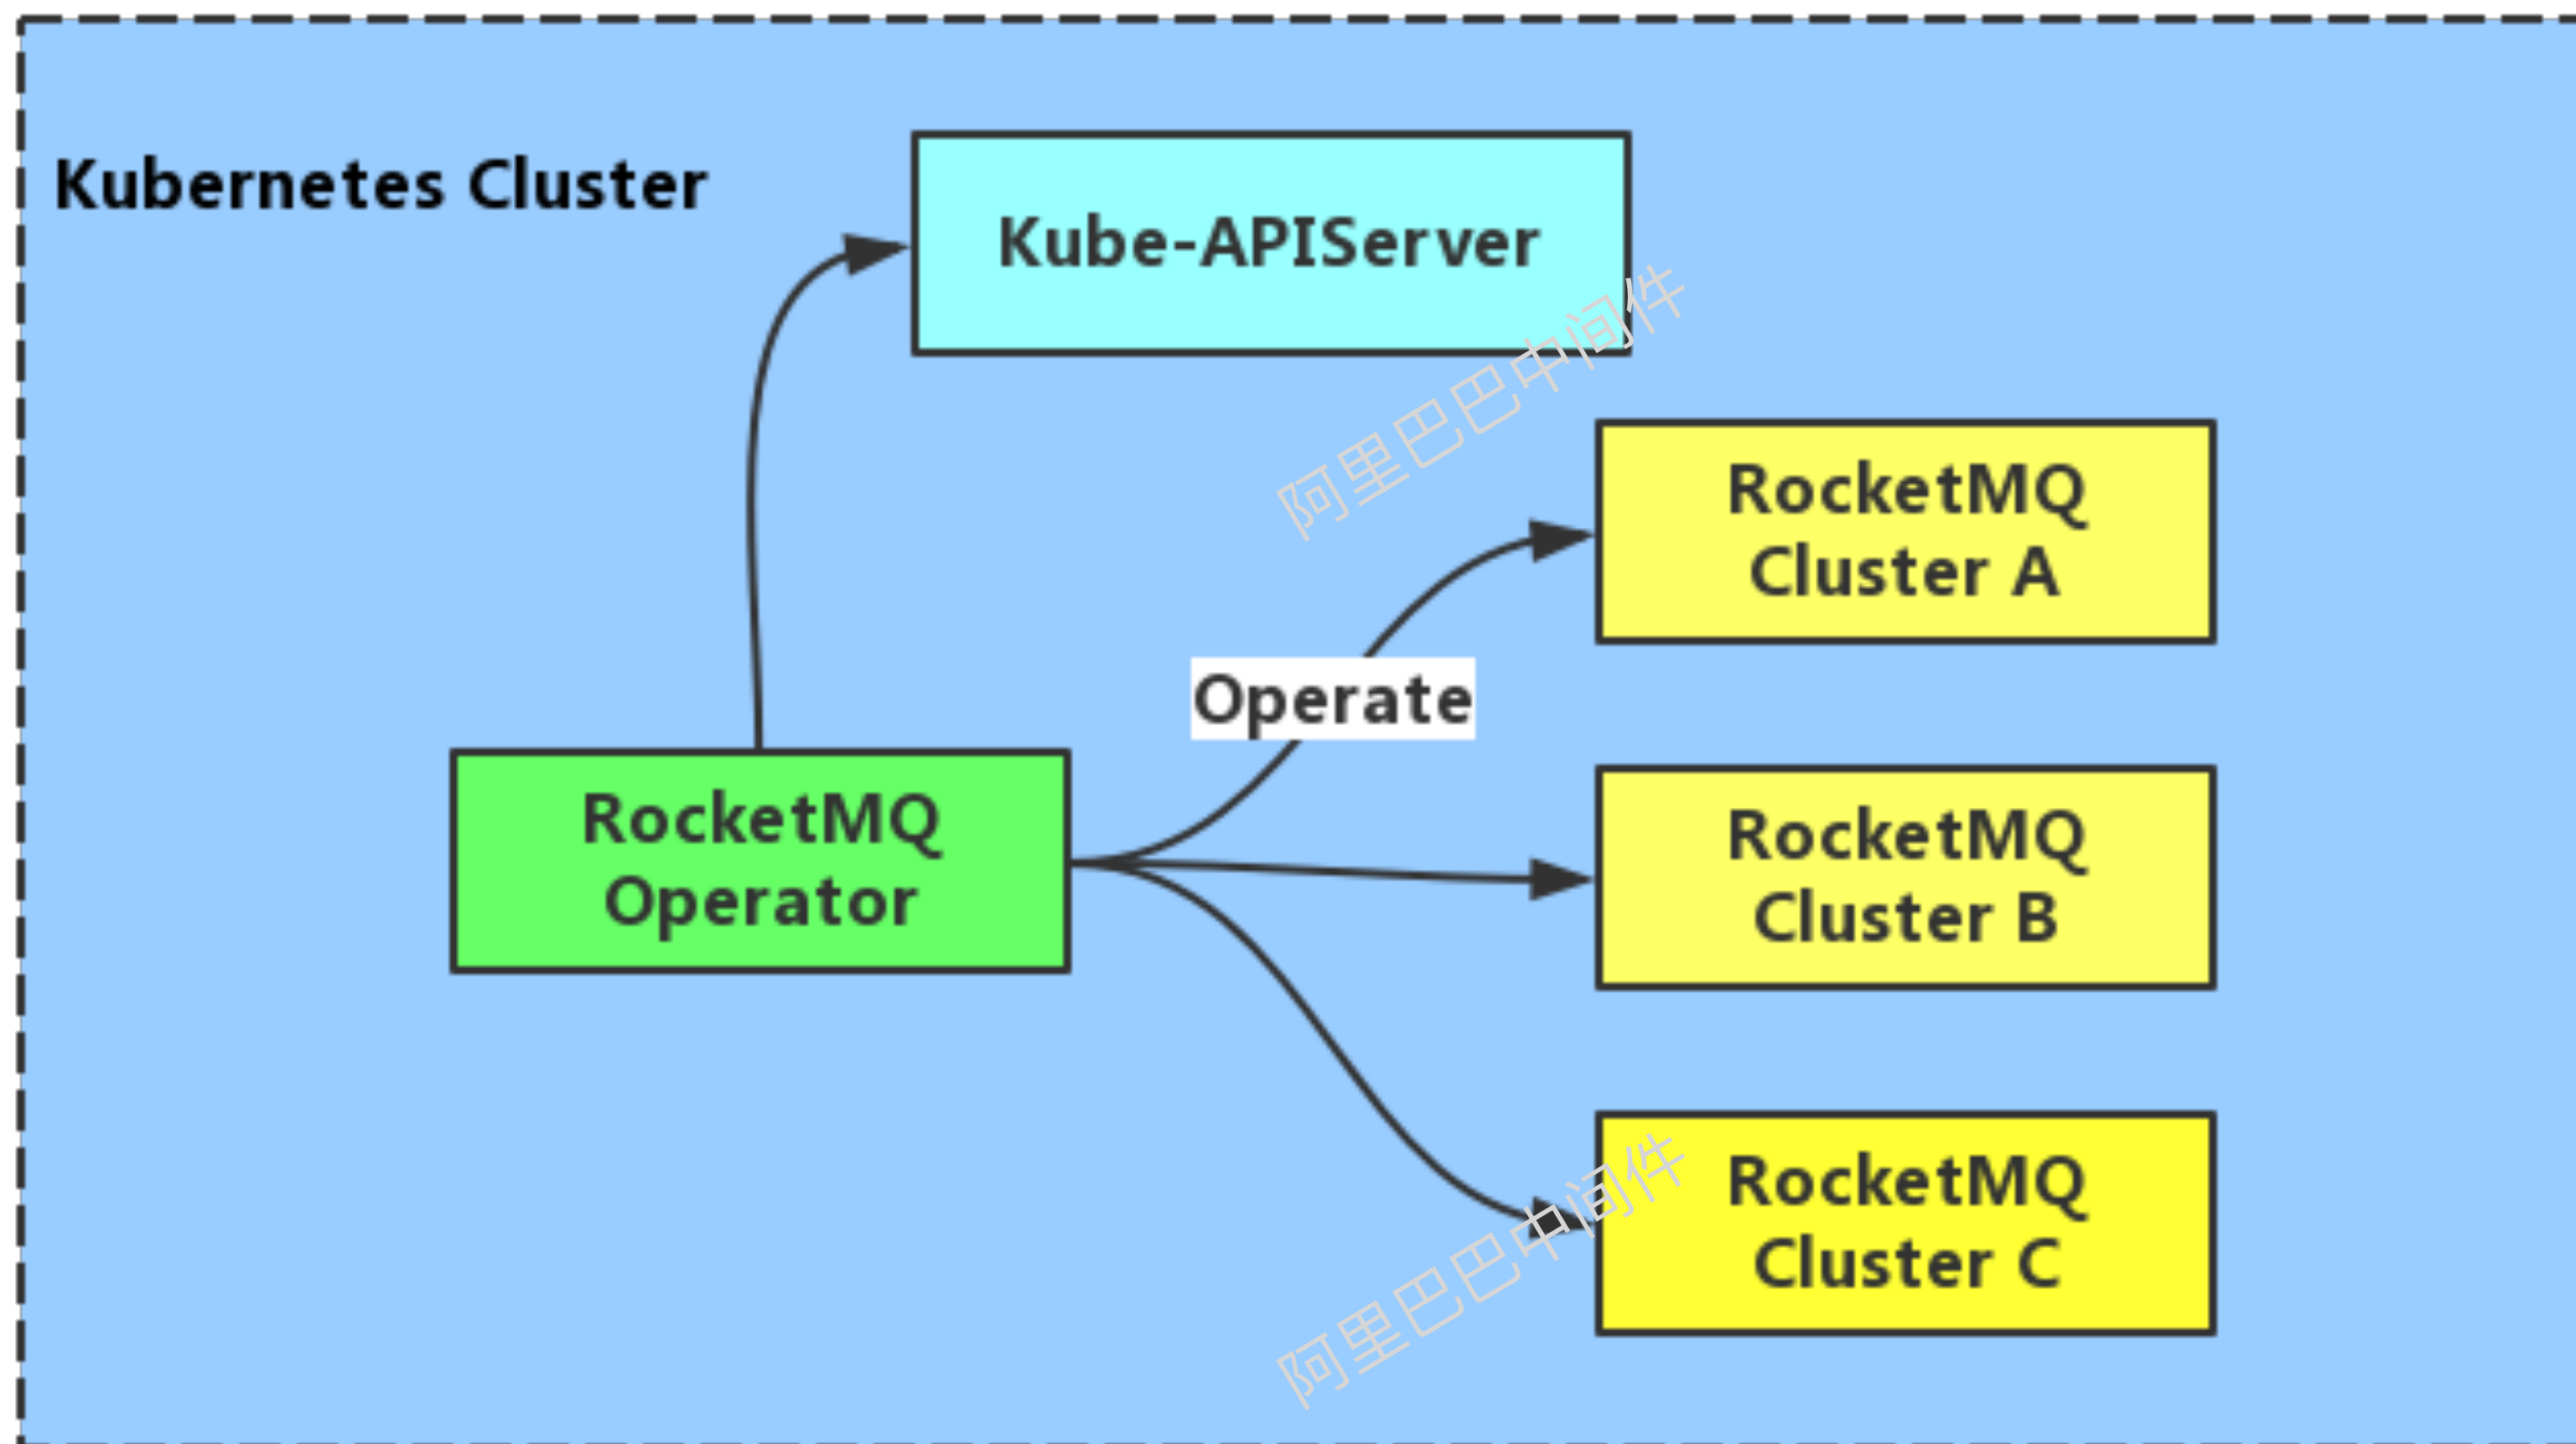
Question: How do we operate RocketMQ clusters?

- How to set up a new RocketMQ cluster in seconds?
- How to scale up/ scale down RocketMQ NameServer endpoints, or Broker endpoints?
- How about RocketMQ endpoints fail over?
- How to upgrade a RocketMQ cluster?
- How to operate all RocketMQ clusters in a data center?
- More concerns ...

What is RocketMQ Operator ?

- It's slogan - Create, operate and scale self-healing RocketMQ clusters on Kubernetes.
- It's a Go program designed for running in Kubernetes cluster to automatically operate RocketMQ clusters.
- It's open sourced. Currently maintained by @huanwei from [HarmonyCloud.cn](https://www.harmonycloud.cn)
- Github - <https://github.com/huanwei/rocketmq-operator>

RocketMQ Operator Topology



Declarative Programming Pattern

```
//just pseudo codes for a Controller  
for {  
    desired := getDesiredState()  
    current := getCurrentState()  
    makeChanges(desired, current)  
}
```

Sample: Controller Pattern In Kubernetes

<https://github.com/kubernetes/kubernetes/blob/release-1.13/cmd/kube-controller-manager/app/controllermanager.go#L357>

```
354
355 // NewControllerInitializers is a public map of named controller groups (you can start more than one in an init func)
356 // paired to their InitFunc. This allows for structured downstream composition and subdivision.
357 func NewControllerInitializers(loopMode ControllerLoopMode) map[string]InitFunc {
358     controllers := map[string]InitFunc{}
359     controllers["endpoint"] = startEndpointController
360     controllers["replicationcontroller"] = startReplicationController
361     controllers["podgc"] = startPodGCController
362     controllers["resourcequota"] = startResourceQuotaController
363     controllers["namespace"] = startNamespaceController
364     controllers["serviceaccount"] = startServiceAccountController
365     controllers["garbagecollector"] = startGarbageCollectorController
366     controllers["daemonset"] = startDaemonSetController
367     controllers["job"] = startJobController
368     controllers["deployment"] = startDeploymentController
369     controllers["replicaset"] = startReplicaSetController
370     controllers["horizontalpodautoscaling"] = startHPAController
371     controllers["disruption"] = startDisruptionController
372     controllers["statefulset"] = startStatefulSetController
373     controllers["cronjob"] = startCronJobController
374     controllers["csrsigning"] = startCSRSigningController
375     controllers["csrapproving"] = startCSRApprovingController
376     controllers["csrcleaner"] = startCSRCleanerController
377     controllers["ttl"] = startTTLController
378     controllers["bootstrapsigner"] = startBootstrapSignerController
379     controllers["tokencleaner"] = startTokenCleanerController
380     controllers["nodeipam"] = startNodeIpamController
381     if loopMode == IncludeCloudLoops {
382         controllers["service"] = startServiceController
383         controllers["route"] = startRouteController
384         // TODO: volume controller into the IncludeCloudLoops only set.
385         // TODO: Separate cluster in cloud check from node lifecycle controller.
386     }
```


RocketMQ Operator Controllers

<https://github.com/huanwei/rocketmq-operator/blob/release-0.1/pkg/controllers/cluster/controller.go#L156>

```
156     m := BrokerController{
157         opConfig: opConfig,
158
159         opClient:  opClient,
160         kubeClient: kubeClient,
161
162         clusterLister:      clusterInformer.Lister(),
163         clusterListerSynced: clusterInformer.Informer().HasSynced,
164         clusterUpdater:      newClusterUpdater(opClient, clusterInformer.Lister()),
165
166         serviceLister:      serviceInformer.Lister(),
167         serviceListerSynced: serviceInformer.Informer().HasSynced,
168         serviceControl:      NewRealServiceControl(kubeClient, serviceInformer.Lister()),
169
170         statefulSetLister:      statefulSetInformer.Lister(),
171         statefulSetListerSynced: statefulSetInformer.Informer().HasSynced,
172         statefulSetControl:      NewRealStatefulSetControl(kubeClient, statefulSetInformer.Lister()),
173
174         podLister:      podInformer.Lister(),
175         podListerSynced: podInformer.Informer().HasSynced,
176         podControl:      NewRealPodControl(kubeClient, podInformer.Lister()),
177
178         queue:      workqueue.NewNamedRateLimitingQueue(workqueue.DefaultControllerRateLimiter(), "brokercluster"),
179         recorder: recorder,
180     }
```

Extending Kubernetes API With Custom Resource Definitions: RocketMQ Operator Types

<https://github.com/huanwei/rocketmq-operator/blob/release-0.1/pkg/apis/rocketmq/v1alpha1/types.go#L22>

```
22  type BrokerClusterSpec struct {
23      BrokerImage      string          `json:"brokerImage"`
24      NameServers       string          `json:"nameServers"`
25      AllMaster         bool           `json:"allMaster"`
26      ReplicationMode    string          `json:"replicationMode"`
27      GroupReplica       int32           `json:"groupReplica, omitempty"`
28      MembersPerGroup    int32           `json:"membersPerGroup, omitempty"`
29      Properties         map[string]string `json:"properties, omitempty"`
30      NodeSelector       map[string]string `json:"nodeSelector, omitempty"`
31      Affinity           *corev1.Affinity  `json:"affinity, omitempty"`
32      VolumeClaimTemplate *corev1.PersistentVolumeClaim `json:"volumeClaimTemplate, omitempty"`
33      Config             *corev1.LocalObjectReference `json:"config, omitempty"`
34  }
35
```

RocketMQ Operator Tutorial

<https://github.com/huanwei/rocketmq-operator/blob/master/docs/tutorial.md>

Prerequisites

- Kubernetes 1.9+, since StatefulSets are stable (GA) in 1.9.

Installation

Create a namespace

```
kubectl create ns rocketmq-operator
```


RocketMQ Operator Tutorial

Deploy the cluster

0) Create a name server cluster

Create a deployment for namesrv, and expose namesrv service with NodePort.

```
kubectl create -f deploy/00-namesrv.yaml
```

1) Create the CRD with kind BrokerCluster.

```
kubectl create -f deploy/01-resources.yaml
```

2) Create RBAC resources for Rocketmq Operator

```
kubectl create -f deploy/02-rbac.yaml
```


RocketMQ Operator Tutorial

3) Deploy Rocketmq Operator

```
kubectl create -f deploy/03-deploymentWithConfig.yaml
```

4) Create a broker cluster

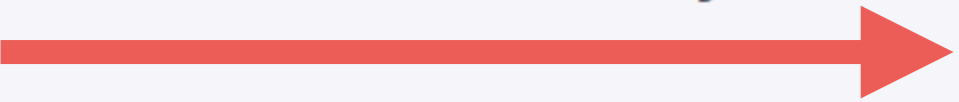
Create a cluster with 2 masters:

```
kubectl create -f deploy/04-cluster-2m.yaml
```

OR:

Create a cluster with 2 masters and 2 slaves:

```
kubectl create -f deploy/04-cluster-2m-2s.yaml
```



```
huan@MacBook-Pro:deploy huan$ cat 04-cluster-2m-2s.yaml
apiVersion: "rocketmq.huanwei.io/v1alpha1"
kind: BrokerCluster
metadata:
  name: mybrokercluster
  namespace : rocketmq-operator
spec:
  brokerImage: huanwei/rocketmq-broker:4.3.0-operator
  nameServers: 10.10.103.181:9876;10.10.101.98:9876
  allMaster: false
  replicationMode: SYNC
  groupReplica: 2
  membersPerGroup: 2
```

RocketMQ Operator Tutorial

Verify the cluster

Below is the command line output for a cluster with 2 masters and 2 slaves.

```
[root@k8s-master deploy]# kubectl get pods -n rocketmq-operator -owide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
mybrokercluster-0-0	1/1	Running	0	38s	192.168.122.145	k8s-node4
mybrokercluster-0-1	1/1	Running	0	33s	192.168.196.17	k8s-node5
mybrokercluster-1-0	1/1	Running	0	38s	192.168.196.16	k8s-node5
mybrokercluster-1-1	1/1	Running	0	34s	192.168.122.146	k8s-node4
rocketmq-namesrv-7d7df8c7d7-4t2g2	1/1	Running	2	10d	10.10.103.181	k8s-node4
rocketmq-namesrv-7d7df8c7d7-6x47m	1/1	Running	2	7d	10.10.101.98	k8s-node5
rocketmq-operator-d9dc87954-nw262	1/1	Running	0	6m	192.168.196.14	k8s-node5

RocketMQ Operator Tutorial

Verify the cluster

Below is the command line output for a cluster with 2 masters and 2 slaves.

```
[root@k8s-master deploy]# kubectl exec -it mybrokercluster-0-0 -n rocketmq-oper
```

```
[root@mybrokercluster-0-0 bin]# cat ../conf/broker.conf
#This file was auto generated by rocketmq-operator. DO NOT EDIT.
brokerClusterName=DefaultCluster
brokerName=broker-0
brokerId=0
deleteWhen=04
fileReservedTime=48
brokerRole=SYNC_MASTER
flushDiskType=ASYNC_FLUSH
```

```
[root@mybrokercluster-0-0 bin]# exit
exit
```

```
[root@k8s-master deploy]# kubectl exec -it mybrokercluster-0-1 -n rocketmq-oper
```

```
[root@mybrokercluster-0-1 bin]# cat ../conf/broker.conf
#This file was auto generated by rocketmq-operator. DO NOT EDIT.
brokerClusterName=DefaultCluster
brokerName=broker-0
brokerId=1
deleteWhen=04
fileReservedTime=48
brokerRole=SLAVE
flushDiskType=ASYNC_FLUSH
```

```
[root@mybrokercluster-0-1 bin]# exit
exit
```

```
[root@k8s-master deploy]# kubectl exec -it mybrokercluster-1-0 -n rocketmq-operator bash
```

```
[root@mybrokercluster-1-0 bin]# cat ../conf/broker.conf
#This file was auto generated by rocketmq-operator. DO NOT EDIT.
brokerClusterName=DefaultCluster
brokerName=broker-1
brokerId=0
deleteWhen=04
fileReservedTime=48
brokerRole=SYNC_MASTER
flushDiskType=ASYNC_FLUSH
```

```
[root@mybrokercluster-1-0 bin]# exit
exit
```

```
[root@k8s-master deploy]# kubectl exec -it mybrokercluster-1-1 -n rocketmq-operator bash
```

```
[root@mybrokercluster-1-1 bin]# cat ../conf/broker.conf
#This file was auto generated by rocketmq-operator. DO NOT EDIT.
brokerClusterName=DefaultCluster
brokerName=broker-1
brokerId=1
deleteWhen=04
fileReservedTime=48
brokerRole=SLAVE
flushDiskType=ASYNC_FLUSH
```

RocketMQ Operator Tutorial

Verify the cluster

Below is the command line output for a cluster with 2 masters and 2 slaves.

```
[root@mybrokercluster-1-1 bin]# ./mqadmin clusterList -n 10.10.103.181:9876 10.10.101.98:9876
```

```
OpenJDK 64-Bit Server VM warning: ignoring option PermSize=128m; support was removed in 8.0
```

```
OpenJDK 64-Bit Server VM warning: ignoring option MaxPermSize=128m; support was removed in 8.0
```

#Cluster Name	#Broker Name	#BID	#Addr	#Version	#InTPS(LC
DefaultCluster	mybrokercluster-0-0	0	192.168.122.145:10911	V4_3_0	0.00(0,0
DefaultCluster	mybrokercluster-0-1	0	192.168.196.17:10911	V4_3_0	0.00(0,0
DefaultCluster	mybrokercluster-1-0	0	192.168.196.16:10911	V4_3_0	0.00(0,0
DefaultCluster	mybrokercluster-1-1	0	192.168.122.146:10911	V4_3_0	0.00(0,0

Project Improvements Plan

- More features are WIP
- Welcome to submit issue/ PR
- Github - <https://github.com/huanwei/rocketmq-operator>

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

阿里巴巴中间件

Thanks!

